

Wine Quality Prediction - AWS Spark Application

Objective:

The goal of this project is to create a Python application using **PySpark** for predicting the quality of wine. The model is trained and tested on an **AWS Elastic MapReduce (EMR)** cluster. Training is parallelized across multiple **EC2 instances**, and the model is deployed using a Docker container for easy scalability and deployment.

Links:

- **GitHub Repository:** <https://github.com/rk94407/rohance>
 - **Docker Hub Repository:** <https://hub.docker.com/repository/docker/rohankatkam1698/testwinequalityprediction/general>
-

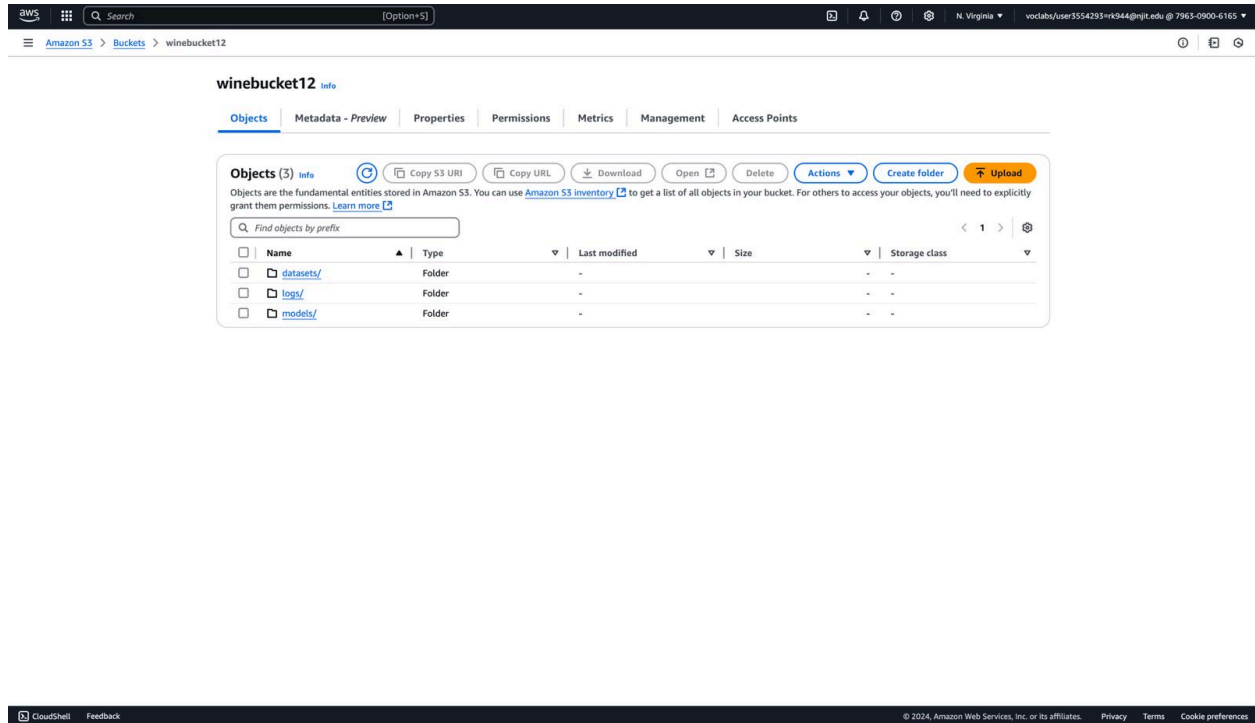
Steps for Execution

1. Create an EC2 Key Pair

- Go to **EC2 > Key Pairs** in your AWS console.
- Generate a new key pair named `hemanth.pem` and download it in `.pem` format for SSH access to the cluster.

2. Set Up an S3 Bucket

- Create an S3 bucket named `winebucket12` to store your datasets and trained models.



3. Launch an EMR Cluster

- Navigate to the **EMR Console** and create a new EMR cluster with the following configurations:
 - **Cluster Name:** winequality
 - **EMR Release Version:** emr-7.5.0
 - **Applications:** Include **Hadoop 3.4.0** and **Spark 3.5.2**.

Clone "winequality"

Name and applications - required

Name your cluster and choose the applications that you want to install to your cluster.

Name: winequality

Amazon EMR release: emr-7.5.0

Application bundle: Spark Interactive

Cluster configuration - required

Choose a configuration method for the primary, core, and task node groups for your cluster.

☒ Uniform instance groups

Summary

Name: winequality

Amazon EMR release: emr-7.5.0

Application bundle: Spark Interactive (Hadoop 3.4.0, Hive 3.1.3, JupyterEnterpriseGateway 2.6.0, Livy 0.8.0, Spark 3.5...)

Cluster configuration - required

Uniform instance groups: Primary (m5.xlarge), Core (m5.xlarge)

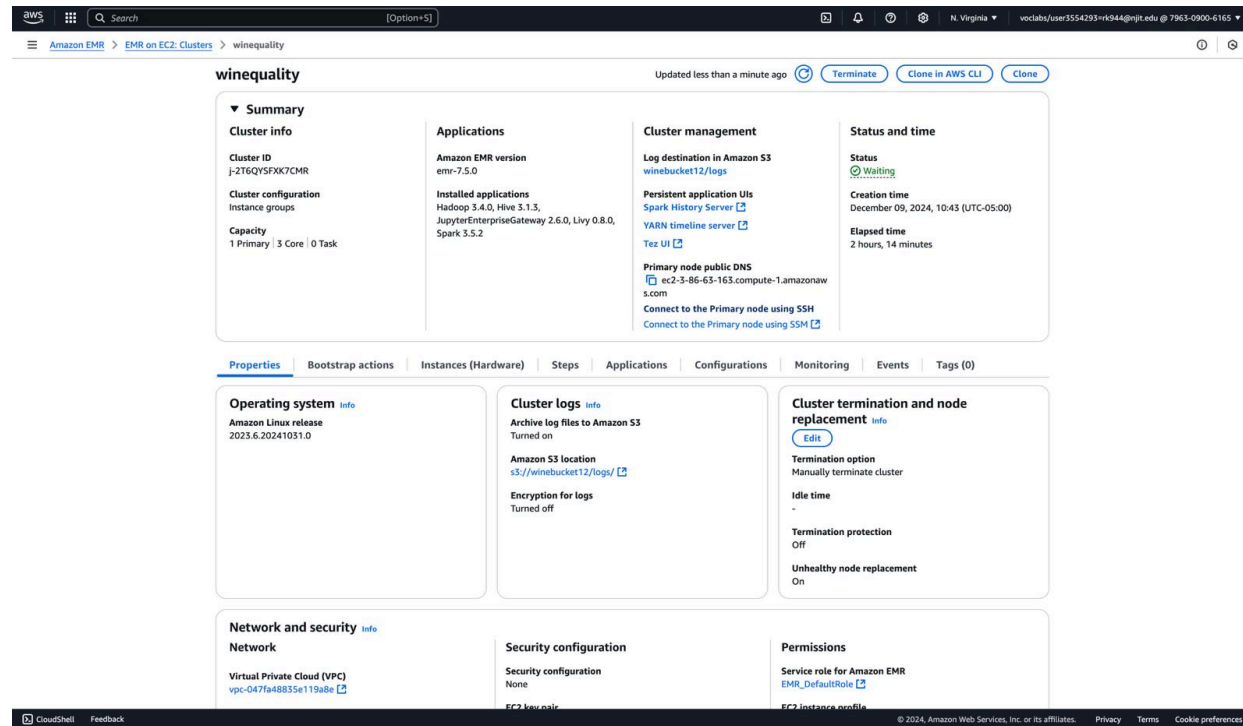
Cluster scaling and provisioning - required

Provisioning configuration: Core size: 3 instances

4. Configure the Spark Cluster

- Use an existing cluster configuration or create a new one.
- Configuration includes:
 - Cluster scaling and provisioning.
 - Networking settings and termination policies.

- Set up **IAM roles** and configure the EC2 key pair (rohan.pem) for security.



5. Train the ML Model on EC2 Instances

Without Docker:

SSH into the **Master Node** of your EMR cluster:

```
ssh -i "rohan.pem" ec2-user@<ec2-public-dns>
```


- This script splits the **TrainingDataset.csv** into 90% for training and 10% for testing. The test data is saved as **TestDataset.csv** in the S3 dataset folder.

6. Save the Trained Model

After training, the model is saved in the following S3 location:

s3://winebucket12/models/winemodel

7. Test the Model

To test the trained model, submit the test job with the following command:

```
spark-submit winequality.py s3://winebucket12/datasets/TestDataset.csv
s3://winebucket12/models/winemodel
```

```
24/12/09 16:59:26 INFO ServerInfo: Adding filter to /stages/pool/json: org.apache.hadoop.yarn.server.webproxy.amfilter.AmIpFilter
24/12/09 16:59:26 INFO ServerInfo: Adding filter to /storage: org.apache.hadoop.yarn.server.webproxy.amfilter.AmIpFilter
24/12/09 16:59:26 INFO ServerInfo: Adding filter to /storage/json: org.apache.hadoop.yarn.server.webproxy.amfilter.AmIpFilter
24/12/09 16:59:26 INFO ServerInfo: Adding filter to /storage/rdd: org.apache.hadoop.yarn.server.webproxy.amfilter.AmIpFilter
24/12/09 16:59:26 INFO ServerInfo: Adding filter to /storage/rdd/json: org.apache.hadoop.yarn.server.webproxy.amfilter.AmIpFilter
24/12/09 16:59:26 INFO ServerInfo: Adding filter to /environment: org.apache.hadoop.yarn.server.webproxy.amfilter.AmIpFilter
24/12/09 16:59:26 INFO ServerInfo: Adding filter to /environment/json: org.apache.hadoop.yarn.server.webproxy.amfilter.AmIpFilter
24/12/09 16:59:26 INFO ServerInfo: Adding filter to /executors: org.apache.hadoop.yarn.server.webproxy.amfilter.AmIpFilter
24/12/09 16:59:26 INFO ServerInfo: Adding filter to /executors/json: org.apache.hadoop.yarn.server.webproxy.amfilter.AmIpFilter
24/12/09 16:59:26 INFO ServerInfo: Adding filter to /executors/threadDump: org.apache.hadoop.yarn.server.webproxy.amfilter.AmIpFilter
24/12/09 16:59:26 INFO ServerInfo: Adding filter to /executors/threadDump/json: org.apache.hadoop.yarn.server.webproxy.amfilter.AmIpFilter
24/12/09 16:59:26 INFO ServerInfo: Adding filter to /executors/heapHistogram: org.apache.hadoop.yarn.server.webproxy.amfilter.AmIpFilter
24/12/09 16:59:26 INFO ServerInfo: Adding filter to /executors/heapHistogram/json: org.apache.hadoop.yarn.server.webproxy.amfilter.AmIpFilter
24/12/09 16:59:26 INFO ServerInfo: Adding filter to /static: org.apache.hadoop.yarn.server.webproxy.amfilter.AmIpFilter
24/12/09 16:59:26 INFO ServerInfo: Adding filter to /: org.apache.hadoop.yarn.server.webproxy.amfilter.AmIpFilter
24/12/09 16:59:26 INFO ServerInfo: Adding filter to /api: org.apache.hadoop.yarn.server.webproxy.amfilter.AmIpFilter
24/12/09 16:59:26 INFO ServerInfo: Adding filter to /jobs/job/kill: org.apache.hadoop.yarn.server.webproxy.amfilter.AmIpFilter
24/12/09 16:59:26 INFO ServerInfo: Adding filter to /stages/stage/kill: org.apache.hadoop.yarn.server.webproxy.amfilter.AmIpFilter
24/12/09 16:59:26 INFO YarnSchedulerBackend$YarnSchedulerEndpoint: ApplicationMaster registered as NettyRpcEndpointRef(spark-client://YarnAM)
24/12/09 16:59:26 INFO ServerInfo: Adding filter to /metrics/json: org.apache.hadoop.yarn.server.webproxy.amfilter.AmIpFilter
24/12/09 16:59:26 INFO YarnClientSchedulerBackend: SchedulerBackend is ready for scheduling beginning after reached minRegisteredResourcesRatio: 0.0
INFO: __main__: Setting Spark log level to ERROR to reduce verbosity...
INFO: __main__: Loading test data from s3://winebucket12/datasets/TestDataset.csv...
INFO: __main__: Cleaning data... Converting all columns to double type.
INFO: __main__: Loading model from s3://winebucket12/models/winemodel...
INFO: __main__: Making predictions on the test dataset...
INFO: __main__: Test Accuracy: 0.7410714285714286
INFO: __main__: Test F1 Score: 0.7289246124028119
INFO: py4j.clientserver: Closing down clientserver connection
[hadoop@ip-172-31-37-199 ~]$
```

Docker Setup

1. Install Docker

- Create a **Docker** account if you don't already have one.
- Install **Docker** on your local machine.

2. Build the Docker Image

Once your Docker environment is set up, build the Docker image for the application:

```
docker build -t testwinequalityprediction .
```

3. Push and Pull the Docker Image

Tag the Docker image:

```
docker tag testwinequalityprediction rohankatkam1698/testwinequalityprediction
```

Push the image to Docker Hub:

```
docker push rohankatkam1698/testwinequalityprediction
```

Pull the image from Docker Hub:

```
docker pull rohankatkam1698/testwinequalityprediction
```

4. Run the Docker Container

After pulling the image, run the container to make predictions using the trained model:

```
docker run --rm testwinequalityprediction /app/datasets/TestDataset.csv /app/models
```

```
(base) mac:rohanc rohankatkam$ docker run --rm testwinequalityprediction /app/datasets/TestDataset.csv /app/models
INFO: __main__:Starting Spark session...
Setting default log level to "WARN".
To adjust logging level use sc.setLogLevel(newLevel). For SparkR, use setLogLevel(newLevel).
24/12/09 17:51:20 WARN NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
INFO: __main__:Loading test data from /app/datasets/TestDataset.csv...
INFO: __main__:Cleaning data... Converting all columns to double type.
INFO: __main__:Loading model from /app/models...
INFO: __main__:Predicting with the model...
24/12/09 17:51:36 WARN DAGScheduler: Broadcasting large task binary with size 9.7 MiB
INFO: __main__:Test Accuracy: 0.7410714285714286
INFO: py4j.clientserver:Closing down clientserver connection
(base) mac:rohanc rohankatkam$
```

Model Accuracy

The **wine quality prediction model** achieved an accuracy of **0.74** based on the test dataset.